

# **CSE 120**

# **Principles of Operating Systems**

**Fall 2014**

**Lecture 12: File Systems**

Geoffrey M. Voelker

# File Systems

---

- First we'll discuss properties of physical disks
  - ◆ Structure
  - ◆ Performance
  - ◆ Scheduling
- Then we'll discuss how we build file systems on them
  - ◆ Files
  - ◆ Directories
  - ◆ Sharing
  - ◆ Protection
  - ◆ File System Layouts
  - ◆ File Buffer Cache
  - ◆ Read Ahead

# Disks and the OS

---


- Disks are messy physical devices:
  - ◆ Errors, bad blocks, missed seeks, etc.
- The job of the OS is to hide this mess from higher level software
  - ◆ Low-level device control (initiate a disk read, etc.)
  - ◆ Higher-level abstractions (files, databases, etc.)
- The OS may provide different levels of disk access to different clients
  - ◆ Physical disk (surface, cylinder, sector)
  - ◆ Logical disk (disk block #)
  - ◆ Logical file (file block, record, or byte #)

# Physical Disk Structure

---

- Disk components

- ◆ Platters
- ◆ Surfaces
- ◆ Tracks
- ◆ Sectors
- ◆ Cylinders
- ◆ Arm
- ◆ Heads



# Disk Interaction

---

- Specifying disk requests requires a lot of info:
  - ◆ Cylinder #, surface #, track #, sector #, transfer size...
- Older disks required the OS to specify all of this
  - ◆ The OS needed to know all disk parameters
- Modern disks are more complicated
  - ◆ Not all sectors are the same size, sectors are remapped, etc.
- Current disks provide a higher-level interface (SCSI)
  - ◆ The disk exports its data as a logical array of blocks [0...N]
    - » Disk maps logical blocks to cylinder/surface/track/sector
  - ◆ Only need to specify the logical block # to read/write
  - ◆ But now the disk parameters are hidden from the OS

# Disks: 2014

---

- Seagate Enterprise Performance 3.5" ([server](#))
  - ◆ capacity: 600 GB
  - ◆ rotational speed: 15,000 RPM
  - ◆ sequential read performance: 233 MB/s (outer) – 160 MB/s (inner)
  - ◆ seek time (average): 2.0 ms
- Seagate Barracuda 3.5" ([workstation](#))
  - ◆ capacity: 3000 GB
  - ◆ rotational speed: 7,200 RPM
  - ◆ sequential read performance: 210 MB/s - 156 MB/s (inner)
  - ◆ seek time (average): 8.5 ms
- Seagate Savvio 2.5" ([smaller form factor](#))
  - ◆ capacity: 2000 GB
  - ◆ rotational speed: 7,200 RPM
  - ◆ sequential read performance: 135 MB/s (outer) - ? MB/s (inner)
  - ◆ seek time (average): 11 ms

# Disk Performance

---

- Disk request performance depends upon three steps
  - ◆ Seek – moving the disk arm to the correct cylinder
    - » Depends on how fast disk arm can move (increasing very slowly)
  - ◆ Rotation – waiting for the sector to rotate under the head
    - » Depends on rotation rate of disk (increasing, but slowly)
  - ◆ Transfer – transferring data from surface into disk controller electronics, sending it back to the host
    - » Depends on density (increasing quickly)
- When the OS uses the disk, it tries to minimize the cost of all of these steps
  - ◆ Particularly seeks and rotation

# Solid State Disks

---

- SSDs are a relatively new storage technology
  - ◆ Memory that does not require power to remember state
- No physical moving parts → faster than hard disks
  - ◆ No seek and no rotation overhead
  - ◆ But...more expensive, not as much capacity
- Generally speaking, file systems have remained unchanged when using SSDs
  - ◆ Some optimizations no longer necessary (e.g., layout policies, disk head scheduling), but basically leave FS code as is
  - ◆ Most commonly, SSDs have the same disk interface (SATA)
  - ◆ But can be faster to use directly over the I/O bus (PCIe)



# Disk Scheduling

---

- Because seeks are so expensive (milliseconds!), the OS tries to schedule disk requests that are queued waiting for the disk
  - ◆ FCFS (do nothing)
    - » Reasonable when load is low
    - » Long waiting times for long request queues
  - ◆ SSTF (shortest seek time first)
    - » Minimize arm movement (seek time), maximize request rate
    - » Favors middle blocks
  - ◆ SCAN (elevator)
    - » Service requests in one direction until done, then reverse
  - ◆ C-SCAN
    - » Like SCAN, but only go in one direction (typewriter)

# Disk Scheduling (2)

---

- In general, unless there are request queues, disk scheduling does not have much impact
  - ◆ Important for servers, less so for PCs
- Modern disks often do the disk scheduling themselves
  - ◆ Disks know their layout better than OS, can optimize better
  - ◆ Ignores, undoes any scheduling done by OS

# File Systems

---

- File systems
  - ◆ Implement an abstraction ([files](#)) for secondary storage
  - ◆ Organize files logically ([directories](#))
  - ◆ Permit sharing of data between processes, people, and machines
  - ◆ Protect data from unwanted access (security)

# Files

---

- A file is data with some properties
  - ◆ Contents, size, owner, last read/write time, protection, etc.
- A file can also have a type
  - ◆ Understood by the file system
    - » Block, character, device, portal, link, etc.
  - ◆ Understood by other parts of the OS or runtime libraries
    - » Executable, dll, source, object, text, etc.
- A file's type can be encoded in its name or contents
  - ◆ Windows encodes type in name
    - » .com, .exe, .bat, .dll, .jpg, etc.
  - ◆ Unix encodes type in contents
    - » Magic numbers, initial characters (e.g., #! for shell scripts)

# Basic File Operations

---

## Unix

- `creat(name)`
- `open(name, how)`
- `read(fd, buf, len)`
- `write(fd, buf, len)`
- `sync(fd)`
- `seek(fd, pos)`
- `close(fd)`
- `unlink(name)`

## NT

- `CreateFile(name, CREATE)`
- `CreateFile(name, OPEN)`
- `ReadFile(handle, ...)`
- `WriteFile(handle, ...)`
- `FlushFileBuffers(handle, ...)`
- `SetFilePointer(handle, ...)`
- `CloseHandle(handle, ...)`
- `DeleteFile(name)`
- `CopyFile(name)`
- `MoveFile(name)`

# File Access Methods

---

- Some file systems provide different **access methods** that specify different ways for accessing data in a file
  - ◆ **Sequential access** – read bytes one at a time, in order
  - ◆ **Direct access** – random access given block/byte number
  - ◆ **Record access** – file is array of fixed- or variable-length records, read/written sequentially or randomly by record #
  - ◆ **Indexed access** – file system contains an index to a particular field of each record in a file, reads specify a value for that field and the system finds the record via the index (DBs)
- **What file access method does Unix, NT provide?**
- Older systems provide the more complicated methods

# Directories

---

- Directories serve two purposes
  - ◆ For users, they provide a structured way to organize files
  - ◆ For the file system, they provide a convenient naming interface that allows the implementation to separate logical file organization from physical file placement on the disk
- Most file systems support multi-level directories
  - ◆ Naming hierarchies (`/`, `/usr`, `/usr/local/`, ...)
- Most file systems support the notion of a current directory
  - ◆ Relative names specified with respect to current directory
  - ◆ Absolute names start from the root of directory tree

# Directory Internals

---

- A directory is a list of entries
  - ◆ <name, location>
  - ◆ Name is just the name of the file or directory
  - ◆ Location depends upon how file is represented on disk
- List is usually unordered (effectively random)
  - ◆ Entries usually sorted by program that reads directory
- Directories typically stored in files
  - ◆ Only need to manage one kind of secondary storage unit



# Basic Directory Operations

---

## Unix

- Directories implemented in files
  - ◆ Use file ops to create dirs
- C runtime library provides a higher-level abstraction for reading directories
  - ◆ opendir(name)
  - ◆ readdir(DIR)
  - ◆ seekdir(DIR)
  - ◆ closedir(DIR)

## NT

- Explicit dir operations
  - ◆ CreateDirectory(name)
  - ◆ RemoveDirectory(name)
- Very different method for reading directory entries
  - ◆ FindFirstFile(pattern)
  - ◆ FindNextFile()

# Path Name Translation

---

- Let's say you want to open `"/one/two/three"`
- What does the file system do?
  - ◆ Open directory `"/` ([well known, can always find](#))
  - ◆ Search for the entry `"one"`, get location of `"one"` (in dir entry)
  - ◆ Open directory `"one"`, search for `"two"`, get location of `"two"`
  - ◆ Open directory `"two"`, search for `"three"`, get location of `"three"`
  - ◆ Open file `"three"`
- Systems spend a lot of time walking directory paths
  - ◆ This is why open is separate from read/write
  - ◆ OS will cache prefix lookups for performance
    - » `/a/b`, `/a/bb`, `/a/bbb`, etc., all share `"/a"` prefix

# File Sharing

---

- File sharing has been around since timesharing
  - ◆ Easy to do on a single machine
  - ◆ PCs, workstations, and networks get us there (mostly)
- File sharing is incredibly important for getting work done
  - ◆ Basis for communication and synchronization
- Two key issues when sharing files
  - ◆ Semantics of concurrent access
    - » What happens when one process reads while another writes?
    - » What happens when two processes open a file for writing?
    - » What are we going to use to coordinate?
  - ◆ Protection

# Protection

---

- File systems implement some kind of protection system
  - ◆ Who can access a file
  - ◆ How they can access it
- More generally...
  - ◆ Objects are “what”, subjects are “who”, actions are “how”
- A protection system dictates whether a given **action** performed by a given **subject** on a given **object** should be allowed
  - ◆ You can read and/or write your files, but others cannot
  - ◆ You can read “/etc/motd”, but you cannot write it

# Representing Protection

## Access Control Lists (ACL)

- For each object, maintain a list of subjects and their permitted actions

## Capabilities

- For each subject, maintain a list of objects and their permitted actions

**Subjects**

	<b>Objects</b>		
	/one	/two	/three
Alice	rw	-	rw
Bob	w	-	r
Charlie	w	r	rw

**ACL**

**Capability**

# ACLs and Capabilities

---

- The approaches differ only in how the table is represented
  - ◆ What approach does Unix use in the FS?
- Capabilities are easier to transfer
  - ◆ They are like keys, can handoff, does not depend on subject
- In practice, ACLs are easier to manage
  - ◆ Object-centric, easy to grant, revoke
  - ◆ To revoke capabilities, have to keep track of all subjects that have the capability – a challenging problem
- ACLs have a problem when objects are heavily shared
  - ◆ The ACLs become very large
  - ◆ Use groups (e.g., Unix)

# File System Layout

---

How do file systems use the disk to store files?

- File systems define a block size (e.g., 4KB)
  - ◆ Disk space is allocated in granularity of blocks
- A “Master Block” determines location of root directory
  - ◆ Always at a well-known disk location
  - ◆ Often replicated across disk for reliability
- A free map determines which blocks are free, allocated
  - ◆ Usually a bitmap, one bit per block on the disk
  - ◆ Also stored on disk, cached in memory for performance
- Remaining disk blocks used to store files (and dirs)
  - ◆ There are many ways to do this

# Disk Layout Strategies


---

- Files span multiple disk blocks
- How do you find all of the blocks for a file?
  1. Contiguous allocation
    - » Like memory
    - » Fast, simplifies directory access
    - » Inflexible, causes fragmentation, needs compaction
  2. Linked structure
    - » Each block points to the next, directory points to the first
    - » Good for sequential access, bad for all others
  3. Indexed structure (indirection, hierarchy)
    - » An “index block” contains pointers to many other blocks
    - » Handles random better, still good for sequential
    - » May need multiple index blocks (linked together)



# Unix Inodes

- Unix inodes implement an indexed structure for files
  - ◆ Also store metadata info (protection, timestamps, length, ref count...)
- Each inode contains 15 block pointers
  - ◆ First 12 are direct blocks (e.g., 4 KB blocks)
  - ◆ Then single, double, and triple indirect



# Unix Inodes and Path Search

---

- Unix inodes are **not** directories
- Inodes describe where on the disk the blocks for a file are placed
  - ◆ Directories are files, so inodes also describe where the blocks for directories are placed on the disk
- Directory entries map file names to inodes
  - ◆ To open “/one”, use Master Block to find inode for “/” on disk
  - ◆ Open “/”, look for entry for “one”
  - ◆ This entry gives the disk block number for the inode for “one”
  - ◆ Read the inode for “one” into memory
  - ◆ The inode says where first data block is on disk
  - ◆ Read that block into memory to access the data in the file

# File Buffer Cache

---

- Applications exhibit significant locality for reading and writing files
- Idea: Cache file blocks in memory to capture locality
  - ◆ Called the **file buffer cache**
  - ◆ Cache is system wide, used and shared by all processes
  - ◆ Reading from the cache makes a disk perform like memory
  - ◆ Even a small cache can be very effective
- Issues
  - ◆ The file buffer cache competes with VM (tradeoff here)
  - ◆ Like VM, it has limited size
  - ◆ Need replacement algorithms again (LRU usually used)

# Caching Writes

---

- On a write, some applications assume that data makes it through the buffer cache and onto the disk
  - ◆ As a result, writes are often slow even with caching
- OSes typically do write back caching
  - ◆ Maintain a queue of uncommitted blocks
  - ◆ Periodically flush the queue to disk (30 second threshold)
  - ◆ If blocks changed many times in 30 secs, only need one I/O
  - ◆ If blocks deleted before 30 secs (e.g., /tmp), no I/Os needed
- **Unreliable, but practical**
  - ◆ On a crash, all writes within last 30 secs are lost
  - ◆ **Modern OSes do this by default; too slow otherwise**
  - ◆ System calls (Unix: fsync) enable apps to force data to disk

# Read Ahead

---

- Many file systems implement “read ahead”
  - ◆ FS predicts that the process will request next block
  - ◆ FS goes ahead and requests it from the disk
  - ◆ This can happen while the process is computing on previous block
    - » Overlap I/O with execution
  - ◆ When the process requests block, it will be in cache
  - ◆ Compliments the disk cache, which also is doing read ahead
- For sequentially accessed files can be a big win
  - ◆ Unless blocks for the file are scattered across the disk
  - ◆ File systems try to prevent that, though (during allocation)

# Summary

---

- Files
  - ◆ Operations, access methods
- Directories
  - ◆ Operations, using directories to do path searches
- Sharing
- Protection
  - ◆ ACLs vs. capabilities
- File System Layouts
  - ◆ Unix inodes
- File Buffer Cache
  - ◆ Strategies for handling writes
- Read Ahead