CSE 120
Principles of Operating Systems

Fall 2014

Lecture 11: Page Replacement

Geoffrey M. Voelker

Memory Management

Final lecture on memory management:

- Goals of memory management
  - To provide a convenient abstraction for programming
  - To allocate scarce memory resources among competing processes to maximize performance with minimal overhead

- Mechanisms
  - Physical and virtual addressing (1)
  - Techniques: Partitioning, paging, segmentation (1)
  - Page table management, TLBs, VM tricks (2)

- Policies
  - Page replacement algorithms (3)

Lecture Overview

- Review paging and page replacement
- Survey page replacement algorithms
- Discuss local vs. global replacement
- Discuss thrashing

Locality

- All paging schemes depend on locality
  - Processes reference pages in localized patterns

- Temporal locality
  - Locations referenced recently likely to be referenced again

- Spatial locality
  - Locations near recently referenced locations are likely to be referenced soon

- Although the cost of paging is high, if it is infrequent enough it is acceptable
  - Processes usually exhibit both kinds of locality during their execution, making paging practical

Demand Paging (OS)

- Recall demand paging from the OS perspective:
    - Pages are evicted to disk when memory is full
    - Pages loaded from disk when referenced again
    - References to evicted pages cause a TLB miss
        - » PTE was invalid, causes fault
    - OS allocates a page frame, reads page from disk
    - When I/O completes, the OS fills in PTE, marks it valid, and restarts faulting process

- Dirty vs. clean pages
    - Actually, only dirty pages (modified) need to be written to disk
    - Clean pages do not – but you need to know where on disk to read them from again

Demand Paging (Process)

- Demand paging is also used when a process first starts up

- When a process is created, it has
  - A brand new page table with all valid bits off
  - No pages in memory

- When the process starts executing
  - Instructions fault on code and data pages
  - Faulting stops when all necessary code and data pages are in memory
  - Only code and data needed by a process needs to be loaded
  - This, of course, changes over time…

Page Replacement

- When a page fault occurs, the OS loads the faulted page from disk into a page frame of memory
- At some point, the process has used all of the page frames it is allowed to use
  - This is likely (much) less than all of available memory
- When this happens, the OS must replace a page for each page faulted in
  - It must evict a page to free up a page frame
- The page replacement algorithm determines how this is done
  - And they come in all shapes and sizes

Evicting the Best Page

- The goal of the replacement algorithm is to reduce the fault rate by selecting the best victim page to remove

- The best page to evict is the one never touched again
  - Will never fault on it

- Never is a long time, so picking the page closest to "never" is the next best thing
  - Evicting the page that won't be used for the longest period of time minimizes the number of page faults
  - Proved by Belady

- We're now going to survey various replacement algorithms, starting with Belady's

Belady's Algorithm

- Belady's algorithm is known as the optimal page replacement algorithm because it has the lowest fault rate for any page reference stream
  - Idea: Replace the page that will not be used for the longest time in the future
  - Problem: Have to predict the future
- Why is Belady's useful then? Use it as a yardstick
  - Compare implementations of page replacement algorithms with the optimal to gauge room for improvement
  - If optimal is not much better, then algorithm is pretty good
  - If optimal is much better, then algorithm could use some work
    » Random replacement is often the lower bound

First-In First-Out (FIFO)

- FIFO is an obvious algorithm and simple to implement
  - Maintain a list of pages in order in which they were paged in
  - On replacement, evict the one brought in longest time ago
- Why might this be good?
  - Maybe the one brought in the longest ago is not being used
- Why might this be bad?
  - Then again, maybe it's not
  - We don't have any info to say one way or the other
- FIFO suffers from "Belady's Anomaly"
  - The fault rate might actually increase when the algorithm is given more memory (very bad)

Least Recently Used (LRU)

- LRU uses reference information to make a more informed replacement decision

  - Idea: We can't predict the future, but we can make a guess based upon past experience

  - On replacement, evict the page that has not been used for the longest time in the past (Belady's: future)

  - When does LRU do well? When does LRU do poorly?

- Implementation

  - To be perfect, need to time stamp every reference (or maintain a stack) – much too costly

  - So we need to approximate it

Approximating LRU

- LRU approximations use the PTE reference bit
  - Keep a counter for each page
  - At regular intervals, for every page do:
    - » If ref bit = 0, increment counter
    - » If ref bit = 1, zero the counter
    - » Zero the reference bit
  - The counter will contain the number of intervals since the last reference to the page
  - The page with the largest counter is the least recently used
- Some architectures don't have a reference bit
  - Can simulate reference bit using the valid bit to induce faults
 - What happens when we make a page invalid?

LRU Clock (Not Recently Used)

- Not Recently Used (NRU) – Used by Unix
  - Replace page that is "old enough"
  - Arrange all of physical page frames in a big circle (clock)
  - A clock hand is used to select a good LRU candidate
    - Sweep through the pages in circular order like a clock
    - If the ref bit is off, it hasn't been used recently
      - What is the minimum "age" if ref bit is off?
    - If the ref bit is on, turn it off and go to next page
  - Arm moves quickly when pages are needed
  - Low overhead when plenty of memory
  - If memory is large, "accuracy" of information degrades
    - What does it degrade to?
    - One fix: use two hands (leading erase hand, trailing select hand)

Example: gcc Page Replace

Example: Belady's Anomaly

Fixed vs. Variable Space

- In a multiprogramming system, we need a way to allocate memory to competing processes

- Problem: How to determine how much memory to give to each process?
  - Fixed space algorithms
    - » Each process is given a limit of pages it can use
    - » When it reaches the limit, it replaces from its own pages
    - » Local replacement
      - Some processes may do well while others suffer
  - Variable space algorithms
    - » Process' set of pages grows and shrinks dynamically
    - » Global replacement
   - One process can ruin it for the rest

Working Set Model

- A working set of a process is used to model the dynamic locality of its memory usage
  - Defined by Peter Denning in 60s
- Definition
  - WS(t,w) = {pages P such that P was referenced in the time interval (t, t-w)}
  - t – time, w – working set window (measured in page refs)
- A page is in the working set (WS) only if it was referenced in the last w references

Working Set Size

- The working set size is the number of pages in the working set

  - The number of pages referenced in the interval (t, t-w)

- The working set size changes with program locality

  - During periods of poor locality, you reference more pages

  - Within that period of time, the working set size is larger

- Intuitively, want the working set to be the set of pages a process needs in memory to prevent heavy faulting

  - Each process has a parameter w that determines a working set with few faults

  - Denning: Don't run a process unless working set is in memory

Example: gcc Working Set

Working Set Problems

- Problems
  - How do we determine w?
  - How do we know when the working set changes?

- Too hard to answer
  - So, working set is not used in practice as a page replacement algorithm

- However, it is still used as an abstraction
  - The intuition is still valid
  - When people ask, "How much memory does Firefox need?", they are in effect asking for the size of Firefox's working set

Page Fault Frequency (PFF)

- Page Fault Frequency (PFF) is a variable space algorithm that uses a more ad-hoc approach
  - Monitor the fault rate for each process
  - If the fault rate is above a high threshold, give it more memory
    » So that it faults less
    » But not always (FIFO, Belady's Anomaly)
  - If the fault rate is below a low threshold, take away memory
    » Should fault more
    » But not always

- Hard to use PFF to distinguish between changes in locality and changes in size of working set

Thrashing

- Page replacement algorithms avoid thrashing
  - When most of the time is spent by the OS in paging data back and forth from disk
  - No time spent doing useful work (making progress)
  - In this situation, the system is overcommitted
    - No idea which pages should be in memory to reduce faults
    - Could just be that there isn't enough physical memory for all of the processes in the system
    - Ex: Running Windows95 with 4 MB of memory…
  - Possible solutions
    - Swapping – write out all pages of a process
    - Buy more memory

Summary

- Page replacement algorithms
  - Belady's – optimal replacement (minimum # of faults)
  - FIFO – replace page loaded furthest in past
  - LRU – replace page referenced furthest in past
    - Approximate using PTE reference bit
  - LRU Clock – replace page that is "old enough"
  - Working Set – keep the set of pages in memory that has minimal fault rate (the "working set")
  - Page Fault Frequency – grow/shrink page set as a function of fault rate

- Multiprogramming
  - Should a process replace its own page, or that of another?

Next time...

- Read Chapters 11, 13